

Query Universal Interface

Haak Saxberg and Jess Hester

December 9, 2011

Contents

1	Introduction	3
2	Language Overview	4
2.1	Basic Computation	4
2.2	Basic Data Structures	4
2.3	Basic Control Structures	4
2.4	Input/Output	4
2.5	Error Handling	5
2.6	Tool Support	5
2.7	Alternatives to QUI	5
3	Example Programs	6
3.1	Model Definition	6
3.1.1	Step-by-step	6
3.1.2	Subclassing a Model	8
3.1.3	Regarding Fields	8
3.2	Using QUI	9
4	Language Design	10
4.1	Getting, Putting, and Creating	10
4.2	Syntax	10
4.3	Semantic Abstractions/Building Blocks	10
5	Language Implementation	11
5.1	Host Language	11
5.2	Parsing	11
5.3	Execution	11
5.4	Backends	11
5.4.1	MongoDB	11
5.4.2	Pickling	11
6	Evaluation	12
A	Non-Relational Databases	14

Listings

1	filemodel.py – defining a simple model	6
2	Session 1 – demonstrating run-time validation	7
3	Session 2 – using fields with functions	8
4	file_subclasses.py – subclassing from FileModel	8
5	Session 3 – creating an instance with arbitrary attributes	9

List of Figures

1	QUI's exception hierarchy	5
2	FileModel's family tree.	6
3	A family tree for FileModel.size.	9
4	QUI's order of operations	12
5	Relational vs. Non-Relational (source: readwriteweb.com)	14

1 Introduction

The need to store persistent data is a headache for many developers. Often, they encounter this problem early on, and due to time or budget pressure, their solution is necessarily tailored to the current realities of their applications – with a limited view, if any, of their application’s future, scaled-up needs.

This makes transporting an application from one storage engine to another a time-consuming and expensive affair; all too often, the costs associated with the migration prevent the application from migrating for an extended period of time, which generally means that users aren’t getting 100% use.

In the past, this issue was annoying but manageable, since persistent storage was frequently implemented using relational databases. By the mid 90’s, nearly all relational databases were mostly SQL-compliant. Every vendor, of course, had their own flavor of SQL which presented the eternal threat that a critical query wouldn’t run on a new database, but for the most part SQL was SQL.

However, there were issues with this approach. Relational databases are inherently constricting — adding new entries to an already-extant database table is easy, but modifying a table to allow for different entry attributes requires halting all user interaction with the table while the table is modified. For popular applications, this presents a nontrivial amount of downtime for a potentially large number of displeased users (see Appendix A for more discussion of the two structures).

In the modern era, companies are increasingly looking to non-relational databases as storage solutions, because they can be scaled up much more cost-effectively and quickly – a must in today’s data-driven business world. However, non-relational databases don’t yet have a common language like SQL; one database’s SELECT could very well be another’s INSERT.

Enter the Query Universal Interface (QUI).

QUI attempts to ease the transition between non-relational databases, like other ORMs do for relational databases. QUI does not pretend to implement a universal querying syntax, but rather a universal interface through which code can talk to data with minimal, if any, change.

2 Language Overview

2.1 Basic Computation

At its heart, QUI is a translation engine. It converts class-like data into a form that is understandable to a backend's API. This form, obviously, can vary quite a bit between databases (thus, the need for QUI in the first place) — and the method of translation is unfortunately specific to each backend.

An overview, however, is possible. QUI translates its own models (user-defined, of course) into what we internally refer to as an interface, which uses the language of the backend specified by the user; this interface is what allows QUI to talk to the backend. Currently, model instance's interface is only created when necessary (during getting and putting), in order to reduce the memory overhead that QUI entails.

In our tests, the performance gains garnered by storing the interface as an instance attribute are not noticeable, especially as models typically have fewer than 20 attributes at a time. Since interfaces are generally classes in their own right, however, they can double (or more) the size of the instance if they're stored; this cost can quickly add up, if the user were to **get/put** numerous **Models**.

2.2 Basic Data Structures

QUI has two foundational elements: **Models** and **Fields**. These are both abstract classes; their purpose is more to serve as interface specifications for the rest of QUI — any future extensions, even those written by others, are guaranteed to integrate seamlessly (as far as QUI is concerned) as long as they honor the requirements of these two root classes.

There are two other 'core' classes behind QUI—**ModelMixin** and **FieldMixin**, which serve similar purposes to **Model** and **Field**: interface requirements, which any future backend mixins need to implement in order to be considered valid. In general, **ModelMixin**'s subclasses are used to add the functionality that **Model** requires of *its* subclasses; likewise with **FieldMixin**'s subclasses and **Field**'s.

2.3 Basic Control Structures

Users can fine-tune their interactions with their databases in a couple of ways:

1. Trivially, by configuring the packaged **Field** subclasses (using keyword arguments defined as part of their model)
2. Defining their own **Field** subclasses, with custom behaviors
3. Defining their own **FieldMixins**, customizing data requirements of their **Models**
4. Defining their own **ModelMixins**, customizing their interaction with the databases

Apart from these, QUI's control structures are Python's control structures, since QUI is an internal DSL.

We'll see an example of configuring **Fields** later on. Creating custom **Fields**, **FieldMixins**, and **ModelMixins** is outside the scope of this report, but works the same way as subclassing any other new-style Python abstract metaclass; as long as you implement the required methods, you have a "valid" subclass with which QUI can work. Whether or not it works the way you intend is an entirely different matter.

2.4 Input/Output

As input, QUI will accept a valid Python class definition. As output, QUI gives users...an altered (dare we say, improved?) class definition. A QUI model need never use the extra methods provided by QUI; model classes can be instantiated just like any other class. Thus, declaring a class to be a QUI model necessitates zero changes to any existing code besides the class definition.

2.5 Error Handling

QUI defines several custom subclasses of `Exception`, `QUIException`, so that the user can be informed about exactly what QUI thinks has gone wrong, and where, and with what values. QUI implements a few “levels” of exception, insofar as Python supports “levels” of program running. At the “compilation” level, QUI raises `ImproperlyConfigured` exceptions for any settings that are not valid or missing, with appropriate error messages alongside.

Every time a `QUIException`, or one of its derived exceptions, is raised, QUI prints an informative error message, detailing its best guess for what caused the exception. Along side, it prints the exception raised by the raw Python—just in case the guess is wrong.

Figure 1 describes the exceptions provided by QUI, and their relation to each other.

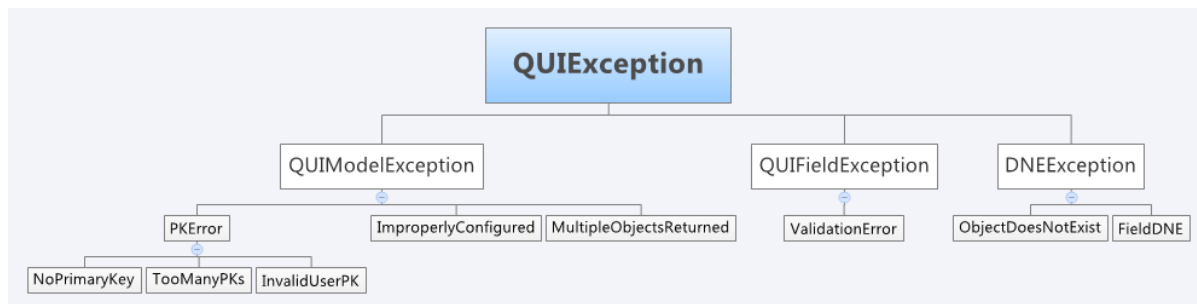


Figure 1: QUI’s exception hierarchy

Although these allow QUI to catch programmatic and procedural errors, we have not implemented an error handling system that intuitively probable *desired* action; if commands are run which are ‘legal’, according to QUI and Python syntax, no errors or warnings are raised—no matter what the effect of those commands are.

2.6 Tool Support

As an internal DSL, QUI “programs” can be written in any environment in which Python programs can be written; the one is a subset of the other. There are no development environments that are *specifically* designed to support or aid QUI programming, but most quality development environments provide class introspection and code completion, which we have found more than sufficient for our own purposes.

Admittedly, we are more familiar with QUI than might be said is average, so your mileage may vary.

2.7 Alternatives to QUI

There are a few interfaces which share a common purpose with QUI – the universalization of database access. None, however, make it their focus to unify *non-relational* databases. Packages like SQLAlchemy and Django’s proprietary ORM work to universalize access to many SQL-based relational databases, but neither officially supports non-relational databases of any sort.

QUI attempts to address an untapped niche market; its declared purpose is unique amongst ORMs.

3 Example Programs

3.1 Model Definition

Here's what it might look like if you wanted to define a simple model:

```

1 from qui.models import Model
2 from qui.fields import *
3 from qui.decorators.storage import stored
4 from qui.decorators.field_decorators import class_field
5
6 @stored(backend="MongoDB")
7 class FileModel(Model):
8     """ A simple file model.
9
10    This model will inherit directly from a supplied mixin - in this case, one
11    that enables the model to talk to MongoDB.
12
13    If your backend changes, the only code you need to change for this model is
14    the the decorator argument - instead of "MongoDB", put "Pickle"
15    (for example), or whatever is appropriate for your backend.
16    """
17
18    count = class_field(IntegerField)
19
20    name = StringField
21    size = IntegerField
22    filetype = StringField
23    notes = StringField
24    created = DateField
25    is_safe = BooleanField
26
27    class_var = 100
28
29    def my_size(self):
30        """ A function specific to this model, unmanaged by QUI """
31        return u"{} bytes".format(self.size)
32
33    def __unicode__(self):
34        if self.name:
35            return u"{}".format(self.name)
36        else:
37            return u"{}".format("Unnamed File")

```

Listing 1: filemodel.py – defining a simple model

FileModel's family tree, after decoration, looks like this:




Figure 2: FileModel's family tree.

3.1.1 Step-by-step

Let's step through this definition. First, we import the necessary things from qui:

```

1 from qui.models import Model
2 from qui.fields import *
3 from qui.decorators.storage import stored
4 from qui.decorators.field_decorators import class_field

```

qui.fields contains all the supported field types, like StringField, IntegerField, etc. At the moment, QUI supports a bare minimum of fields - just the most basic types.

Next, we decorate the class with `storage()`:

```

6 @stored(backend="MongoDB")
7 class FileModel(Model):

```

the `backend` keyword argument to the `storage` decorator tells QUI which mixins it needs to add to your model's inheritance tree, as well as the inheritance trees of any `Fields` you've defined for the model. `storage` will accept arbitrary keyword arguments, but only acts on a few:

1. `backend`, which defaults to `None` and will (generally) raise an `ImproperlyConfigured` exception if not defined, or if you pass in a backend-identifying string that isn't recognized.
2. `host`, which takes a string and defaults to `'localhost'`.
3. `db`, which takes a string and defaults to a backend-specific value.
4. `port`, which takes an integer, and defaults to a backend-specific port.

Our simple `FileModel`, then, uses the MongoDB backend, and will be stored locally, with default settings. Any subclasses of `FileModel` will automatically inherit these settings, but fine control is possible by defining `_host`, `_db`, `_port` on the subclass in question. If, for whatever reason, we decided we'd rather store `FileModels` using Pickling (for example), we'd have to change line 6 to `@stored(backend='Pickle')`, and QUI would acquiesce, no questions asked.

Next, we set some fields that we think a `FileModel` should have:

```

18     count = class_field(IntegerField)
19
20     name = StringField
21     size = IntegerField
22     filetype = StringField
23     notes = StringField
24     created = DateField
25     is_safe = BooleanField
26
27     class_var = 100

```

Notice that we don't *initialize* any of the fields here - we're just defining what kind of `Field` each attribute should be. Although QUI will store any attributes that an instance of `FileModel` may have (except private ones and callable ones), these `Field` attributes will do run-time compatibility checking. For example:

```

1 $ python -i FileModel.py
2 >>> f = FileModel()
3 >>> f.size = "really big"
4 Traceback (most recent call last):
5 ...
6 ...
7 fields.ValidationError: Could not convert to int: really big
8 >>> f.size = 100
9 >>> f.size
10 100
11 >>>

```

Listing 2: Session 1 – demonstrating run-time validation

`count` has been marked as a class attribute by the `class_field` decorator, and it will behave just like any other class attribute. Also, notice that you can still have “regular” class variables, like `class_var`; these will be stored by QUI just like all other non-callable attributes.

There's a (pedantic) function definition, just to show that we can, and (finally) we have `__unicode__`, which controls the way that a QUI model is printed via the `print` commands. :

```

29     def my_size(self):
30         """ A function specific to this model, unmanaged by QUI """
31         return u"{} bytes".format(self.size)
32
33     def __unicode__(self):
34         if self.name:
35             return u"{}".format(self.name)
36         else:
37             return u"{}".format("Unnamed File")

```

Since `my_size` is a callable, QUI ignores it - but doesn't remove it from the model.

```

1 $ python -i FileModel.py
2 >>> f = FileModel()
3 >>> f
4 <FileModel: Unnamed File>
5 >>> print f
6 Unnamed File
7 >>> isinstance(f, FileModel)
8 True
9 >>> f.size
10 >>> f.my_size()
11 'None bytes'
12 >>> f.size = 100
13 >>> f.my_size()
14 '100 bytes'
15 >>>

```

Listing 3: Session 2 – using fields with functions

3.1.2 Subclassing a Model

```

1 from filemodel import FileModel
2 from qui.decorators.storage_decorators import subclass
3 from qui.decorators.field_decorators import class_field
4
5 class ImageFile(FileModel):
6     pass
7
8
9 class TextFile(FileModel):
10     word_count = IntegerField
11     txt_file_count = 100
12
13 @subclass
14 class MoreClassFields(FileModel):
15     special_count = class_field(IntegerField)
16
17 class RemoteFile(FileModel):
18     _port = 91711
19
20     def __init__(self, host):
21         super(Remote_File, self).__init__(self)
22         self._host = host

```

Listing 4: file.subclasses.py – subclassing from FileModel

This is just an example of a few ways that you can customize your QUI models. `ImageFile` is a direct subclass of `FileModel`, and has no further customizations. QUI will store it instances of `ImageFile` in a different ‘place’ than it does instances of `FileModel`, if that makes sense for the backend.

`TextFile` shows that you can define further `Fields` or regular class member variables, as you’d expect from a subclass.

`MoreClassFields` demonstrates the `subclass` decorator, which is only needed if your subclass defines more class-wide `Fields`.

`Remote_File` exhibits the ability to redefine where the model should be stored, by customizing `_host` and `_port` members. The difference between defining them as class members and as instance members is subtle: all QUI functions will use the class member values, if they exist; only `put()` will use the instance values. So, `Remote_Files` will be **gotten** from the wherever `FileModel` has said to find them, but will be **put** to wherever `_host` has been initialized to. In this case, if `x = Remote_File.create('otherhost')`, it will be **created** to `FileModel`’s host (localhost), but on port `91711`. Then `x.put()` will store it on otherhost, also using port `91711`. This feature is great for migrating data between hosts; you’ll be `get()`ing from one host and `put()`ing to another!

Note that if you define a custom `__init__`, you also need to call `super`’s `__init__` before you start modifying any instance attributes, so that QUI can do the necessary setup.

3.1.3 Regarding Fields

We really can’t stress enough that you generally *do not initialize Fields yourself*. All direct subclasses of `Field` are abstract, and even if they weren’t, they do not define all the methods that `Field` requires

in order to be instantiated. They have to be combined with a subclass of `FieldMixin` which fulfills their remaining requirements.




Figure 3: A family tree for `FileModel.size`.

QUI does this for you when you decorate a class definition, and instantiates the resulting fields at the appropriate times—so that every `FileModel` has its own `name`, but the same `count`. Defining `Fields` on your model isn't required; one of the best parts about non-relational databases is that they don't enforce a strict schema. Remember `x`, from above? If we had, somewhere down the line, decided that `x.sensitive = True`, QUI would happily store the `sensitive` attribute right along with all of `x`'s other attributes. Defining an attribute as a `Field` is mostly useful for enforcing a type, or coercing into a certain format:

```

1 >>> x.name = 123
2 >>> x.name
3 u'123'

```

3.2 Using QUI

QUI's ease of use is best exhibited by entering the Python interpreter:

```

1 $ python -i FileModel.py
2 >>> f = FileModel.create(
3 ...   name="arrow.jpg",
4 ...   size=100,
5 ...   title="Greenboy")
6 >>> f
7 <FileModel: arrow.jpg>
8 >>> f.title
9 'Greenboy'
10 >>> f.size
11 100
12 >>> f.size = "really big"
13 Traceback (most recent call last):
14 ...
15 ...
16 fields.ValidationError: Could not convert to int: really big
17 >>>

```

Listing 5: Session 3 – creating an instance with arbitrary attributes

4 Language Design

4.1 Getting, Putting, and Creating

The idea behind QUI is to support a uniform set of ‘verbs’ across all backends. As an internal language focusing on storage, we implemented these verbs as class methods and instance methods.

Since our time for this project was limited, we focused our efforts on what we saw as the three most important verbs for a persistent storage solution:

1. **Creating** an instance of a stored model, at run-time, on the fly, and having it stored into the backend.
2. **Putting** an instance of a stored model into the backend.
3. **Getting** an instance of a stored model from the data stored in the backend.

This tight focus on basic database operations also means there is almost no difficulty in learning QUI’s syntax.

4.2 Syntax

As has been said repeatedly, QUI is an internal DSL; its syntax is no different from that allowed in any valid Python program. Nevertheless, QUI strives (in true Pythonic style) for an optimal combination of clarity and concision. It achieves this by only officially exposing three extra methods on stored models to users: `get(**kwargs)`, `create(**kwargs)`, and `put()`.

By only adding these three verbs to any `Model` definition, QUI runs the least risk of becoming a burden to anyone who decides to use it; furthermore, these method names are unlikely to be already overridden. QUI doesn’t want to step on anybody’s toes, after all.

Most of QUI’s syntax is covered in Section 3; the example `Model` definitions and interpreter sessions cover the fundamentals, and even some non-fundamentals, of using QUI pretty extensively. The user has a large amount of power when decided how exactly to create their Models, and are able to give the Model functionality in whatever manner they deem fit.

4.3 Semantic Abstractions/Building Blocks

QUI is abstracted into a number of different, though related sections. The user creates some sort of model (e.g., a `FileModel`), and performs `put`, `get`, and `create` operations as needed (they can also create other subclasses or functions within the `FileModel` class). At the top of the file (for now), users must import our `Model`, any fields they want to use, and any decorators they want to use.

Under the hood, operations become more complex for each API. At its core, however, QUI uses the ‘black box’ of the API for each backend database, and accesses methods as necessary to create, put, or get model instances. The user inputs commands using their created file (see examples above), which is decorated with the appropriate backend. These decorators call the appropriate Field and Model mixins so attributes are properly assigned and stored depending on the backend. The user should never see any actual database operations; QUI will handle all connections and operations necessary beyond the user’s call to `get`, `put`, or `create`.

5 Language Implementation

5.1 Host Language

QUI is implemented in Python, which was chosen for three main reasons:

1. Our own familiarity with the language.
2. Excellent metaprogramming facilities.
3. Availability of backend APIs.

In a big way, our vision of QUI's design steered us towards Python as the host language.

5.2 Parsing

As an internal DSL, it's unnecessary for QUI to perform any 'parsing' operations, in the same sense as an external one — there's no need for an abstract syntax tree, as QUI syntax is merely Python syntax. In addition, we can rely on Python to catch syntactically illegal expressions in many cases.

5.3 Execution

As an internal DSL, QUI's semantics do not differ from those of its host language, Python. Figure 4 describes the operations that QUI is doing behind the scenes when a class is decorated with `storage`. The `subclass` decorator is similar, but only performs the operations relating to class-wide fields, since a subclass of a QUI model will already have much of the setup done by its parent class.

5.4 Backends

Currently, two backends are supported.

5.4.1 MongoDB

MongoDB is a popular, free non-relational database. For putting, model instances are converted to JSON, and then saved in a database via some port. The port default is 27017, but the user can specify the port via passed in arguments. The `_get_interface()` method takes the model instance and sets each field as a key:value pair in the JSON dict.

Getting is simply a matter of finding the object which corresponds to the passed in keyword arguments.

5.4.2 Pickling

Also supported is local storage of model instances, using the Python pickle module. Model instances are stored in the directory specified by the `_path` variable, in the file `Modelname\InstancePrimaryKey`. Sadly, Python's pickle module will only pickle global classes (namely, not our `Models`), so this backend simply pickles a dict which contains the state of the model instance.

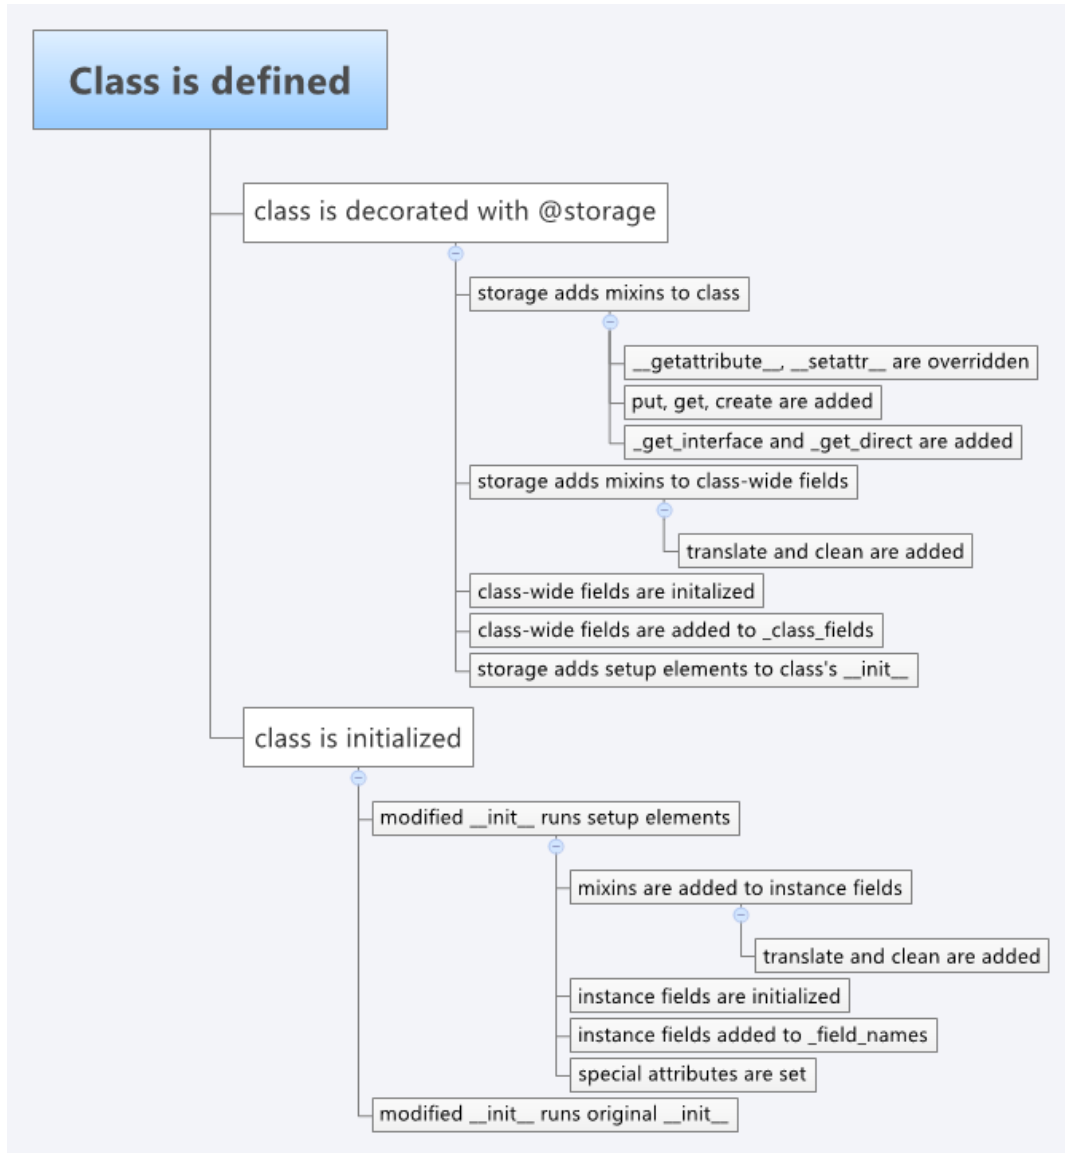


Figure 4: QUI's order of operations

6 Evaluation

The basic implementation of QUI went as planned, building the framework for adding various backends. Adding a new backend is a nontrivial task, but only one file needs to be added (`NewBackEndMixins.py`) defining the necessary classes, and QUI will take care of performing the operations for a `Model` decorated to use that backend.

Likewise, QUI's syntax is clear and concise, providing basic functionality while allowing the user full access to the power of Python. The user is free to use our built-in backend mixins, or to define their own if they so wish. This is well in keeping with the idea of non-relational databases: to give the user the flexibility they need to extend their applications in whatever directions they wish to go. Furthermore, we built in several 'escape hatches,' to allow users to avoid QUI's topmost layer, if desired: methods like `_get_direct()` allow a user to directly access a `Field`, bypassing `Model`'s overridden `__getattrible__` and `__setattr__` entirely.

However, some difficulties were encountered in the implementation of QUI. Our initial intention was to use Google AppEngine as another backend, but upon investigation into the AppEngine datastore, we found that AppEngine supports a non-relational database solely for applications on Google AppEngine. Pulling out just the datastore connection proved essentially impossible without a working application

online. In light of this, providing support for just the AppEngine database seemed both overly time-consuming and unlikely to benefit any users, as they still have to use other AppEngine functionalities to access their datastores. So, we focused on MongoDB and Pickling as our supported backends.

Time constraints also presented an issue during later development. The two methods required for specific backend FieldMixins have not yet been fully implemented, as they have almost no impact on the usability of the database and were added as a potential feature that might be convenient to users. The method `translate()` currently returns the value of a field, but would in the future convert a top-level `StringField()` (for example) to a `db.StringProperty()`, or whatever is appropriate for that backend.

Likewise, the method `clean()` is intended to sanitize data before it is inserted into the database.

QUI was developed and implemented as a team. The initial design was done entirely on whiteboards with both participants working together.

Later in the process, pair programming was used, with both members working on one computer. This was the abstract base classes and later the classes for the mixins to inherit.

The last stage of the project was actually developing for the backends. The backends were split between the two team members, though both had input in each backend — for example, the decision to stop trying to support AppEngine as a backend was made as a team.

Appendices

A Non-Relational Databases

Non-relational databases differ distinctly from their more popular cousins. Relational databases adhere to the more traditional structure of some number of tables, linked by specific keys, with some number of rows and columns. Every row in a specific table must have some specific attributes (the columns of the table), the same as every other row of the table. Adding a column to a table then requires stopping all interactions on that table, as every entry must be changed.

Non-relational databases are not as restricted in their structure (see Figure 5). Rather than each entry being a row in a table with a fixed number of columns, each model instance is a 'page' in the database, with its own fields (some of which may be required, and others specific to the instance itself). So, non-relational databases are much more horizontally flexible when freezing the database would be inopportune if a new field needs to be added.

Implementations of non-relational databases vary widely between backends, but the main idea of non-relational databases is removing the field restriction for the user. One of QUI's basic goals is to preserve this functionality.

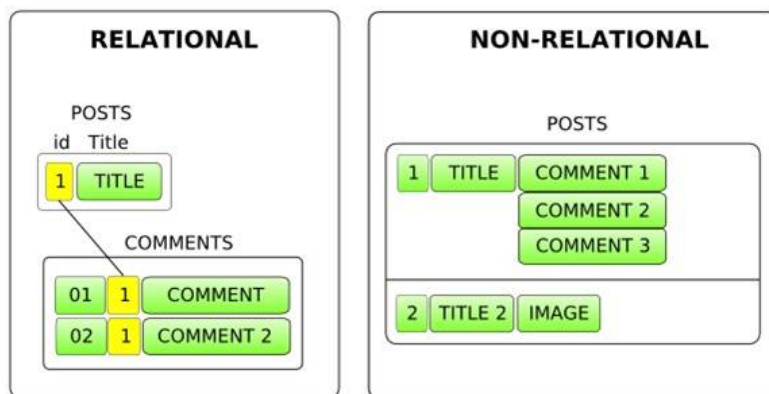


Figure 5: Relational vs. Non-Relational (source: readwriteweb.com)